

Jake Hyun

🏠 jakehyun.com ✉ jakehyun@snu.ac.kr 🎓 Google Scholar 📄 GitHub 🔗 LinkedIn

Research Interests

My interests lie in the intersection of Computer Systems and Artificial Intelligence. I specialize in enhancing computational efficiency and improving model performance, with experience in code parallelization/acceleration across tasks like dimensionality reduction, NLP, and LLM quantization with custom CUDA kernels. Recently, I have been deeply involved in constructing novel LLM inference schemes for edge devices, participating in AI competitions, and developing techniques to push the boundaries of LLM inference speed and accuracy under constrained resources.

Education

2019-2025 **Seoul National University**,
B.S. in Computer Science & Engineering, minor in Linguistics, GPA 3.76/4.3.
Leave of absence for military service: Mar 2021 - Sep 2022

Research Experience

- 2023-2024 **SNU Architecture and Code Optimization Lab**, *Internship*, (Advisor: Prof. Jae W. Lee).
 - Contributed to Any-Precision LLM experimentation and significantly enhanced the core quantization logic, achieving performance gains exceeding 10,000x.
 - Working on a novel inference scheme for quantized LLMs on edge platforms - to be disclosed shortly.
- 2023 **SNU Human-Computer Interaction Lab**, *UROP*, (Advisor: Prof. Jinwook Seo).
 - Optimized Dimensionality Reduction (DR) technique UMATO to reach performance comparable to SOTA competitors.
 - Helped create ZADU, enabling efficient & comprehensive evaluation of DR embeddings through optimized distortion measures.

Publications

- Under Review **UMATO: Bridging Local and Global Structures for Reliable Visual Analytics with Dimensionality Reduction**,
TVCG *Hyeon Jeon, Kwon Ko, Soohyun Lee, Jake Hyun, Taehyun Yang, Gye-hun Go, Jaemin Jo, and Jinwook Seo*.
 - Optimized UMATO, a novel SOTA dimensionality reduction technique that captures both local and global structures.
- ICML Oral **Any-Precision LLM: Low-Cost Deployment of Multiple, Different-Sized LLMs**, [PDF](#)
2024 *Yeonhong Park, Jake Hyun, SangLyul Cho, Bonggeun Sim, Jae W. Lee, Oral Presentation (Top 1.5%)*.
 - Optimized core Any-Precision LLM quantization logic, and worked on model performance experimentations.
 - Implemented the official open-source library, enabling automatic application of our work on arbitrary LLMs.
- IEEE VIS **ZADU: A Python Library for Evaluating the Reliability of Dimensionality Reduction Embeddings**, [PDF](#)
2023 *Hyeon Jeon, Aeri Cho, Jinhwa Jang, Soohyun Lee, Jake Hyun, Hyung-Kwon Ko, Jaemin Jo, Jinwook Seo*.
 - Converted GPU-accelerated evaluation code to optimize performance on CPU platforms.

Awards & Achievements

- 2024 **Accelerator Programming Winter School, CUDA competition**, *1st place, team of two*,
[Organized by SNU THUNDER Research Group & Manycoresoft].
 - 1st place by performance, final project on model inference throughput optimization using CUDA C++.
- 2022 **Korean AI Competition**, *1st place (Undergrad Div.), team of four, prize: KRW 10M*,
[Organized by Korea Ministry of Science and ICT, National Information Society Agency].
 - Developed a speech-to-text model for the Korean language & its dialects.
 - Awarded by Korean Minister of Science and Technology.
- 2020 **SNUH Medical AI Challenge**, *4th place, team of 11*,
[Organized by Seoul National University Hospital].
 - Developed an intraoperative hypotension predictor from arterial pressure waveforms.
- 2020 **Digital Health Hackathon**, *1st place, team of three, prize: KRW 3M*,
[Organized by Samsung Advanced Institute for Health Sciences & Technology, Digital Healthcare Partners].
 - Created a drug treatment decision model for rare cancer, based on a two-model ensemble approach.
- 2017 **Korean Olympiad in Informatics, project division**, *Silver(3rd place)*,
[Organized by Korea Ministry of Science and ICT].
 - Created an RL based AI agent for the game of Othello and Omok.

Open-Source Contributions

- 2024 **Any-Precision LLM**, [repo link](#),
[An LLM quantization library capable of quantizing and running variable bit-width models].
- Implements work from ICML paper of the same name (listed above).
 - Created the highly optimized yet versatile quantization pipeline.
- 2024 **flash1dkmeans**, [repo link](#),
[An optimized K-means implementation for the one-dimensional case].
- Devised, verified and implemented a variant of K-means clustering highly optimized for the 1D data.
 - Used directly in quantization works like Any-Precision LLM to bring down the quantization cost dramatically.
- 2023 **Steadiness & Cohesiveness**, [repo link](#),
[Quality metrics for evaluating the inter-cluster reliability of multidimensional projections].
- Parallelized the algorithm for distance matrix and Shared Nearest Neighbor (SNN) matrix calculations.
- 2023 **UMATO: Uniform Manifold Approximation with Two-phase Optimization**, [repo link](#),
[A dimensionality reduction technique that preserves both global and local structures of high-dimensional data].
- Reduced time complexity of core algorithm to reach performance comparable to UMAP, a SOTA technique.

Academic Project Highlights

- 2023 **LLVM Compiler Optimization Project**, *1st of 11 teams, Principles and Practices of Software Development*.
- Secured 1st place in competitive project optimizing compiler performance on custom system using LLVM passes.
- 2023 **Crowd-Analyzer**, *Creative Integrated Design*.
- Developed a CCTV crowd flow and density analysis system for security platform company INNODEP.
- 2022 **System Programming Lab Assignments**, *1st of 107, System Programming*.
- 1st by optimization score in tasks encompassing IO, dynamic memory, shell design and socket programming.
- 2020 **Computer Architecture Optimization Tasks**, *1st of 130, Computer Architecture*.
- 1st in all 4 competitive assignments, including floating point conversion and RISC-V assembly programming.

Personal Projects

- 2023 **Brick Breaker AI**, [repo link](#).
- Created a clone of a popular brick breaker game variant.
 - Trained a neural network on self-generated labels, bootstrapping the process to create a competent AI agent.
- 2021 **Tranquil Tempest: An ISMCTS based Mighty AI**, [repo link](#).
- Personal project implementing an AI for the card game of Mighty based on ISMCTS, a variant of the MCTS algorithm for imperfect information games.

Extra Curriculars

- 2023 Working as website and server administrator for SNUPO(Seoul National University Philharmonic Orchestra).
- 2021-2022 Served as the Cybersecurity Team at the Cyber Defense Center, 7th Corps of ROK Army.
- 2020 Developed an SMS survey scheduler for Center for Happiness Studies, Seoul National University.
- 2019, 2020 UCPC(Union of Clubs for Programming Contests) programming contest finalist.
Participated in ACM-ICPC 2019 and 2020, Google Code Jam 2019 and 2020.
- 2018 ISEF-K(International Science and Engineering Fair Korea) participant, with improved Othello/Omok AI over the KOI competition(2017).

Skills

- Languages Python, C, C++
- Frameworks PyTorch, Django, Numba
- Technologies CUDA, LLVM Passes
- Proficiencies Arduino, Raspberry Pi, Over 12 years of GNU/Linux experience.

Personal Details

- Language English & Korean, bilingual proficiency.